

Explore language understanding

 **Note** To complete this lab, you will need an [Azure subscription](#) in which you have administrative access.

Increasingly, we expect computers to be able to use AI in order to understand spoken or typed commands in natural language. For example, you might want to implement a home automation system that enables you to control devices in your home by using voice commands such as “switch on the light” or “put the fan on”, and have an AI-powered device understand the command and take appropriate action.

To test the capabilities of the Conversational Language Understanding service, we’ll use a command-line application that runs in the Cloud Shell. The same principles and functionality apply in real-world solutions, such as web sites or phone apps.

Create a *Language service* resource

You can use the Conversational Language Understanding service by creating a **Language service** resource.

If you haven’t already done so, create a **Language service** resource in your Azure subscription.

1. In another browser tab, open the Azure portal at <https://portal.azure.com>, signing in with your Microsoft account.
2. Click the **+ Create a resource** button, search for *Language service*, and create a **Language service** resource with the following settings:
 - Select additional features: *Keep the default features and click Continue to create your resource*
 - **Subscription:** *Your Azure subscription.*
 - **Resource group:** *Select or create a resource group with a unique name.*
 - **Region:** East US 2
 - **Name:** *Enter a unique name.*
 - **Pricing tier:** S (1K Calls per minute)
 - **By checking this box I certify that I have reviewed and acknowledge the terms in the Responsible AI Notice.:** Selected.
3. Review and create the resource, and wait for deployment to complete.

Create a Conversational Language Understanding App

To implement natural language understanding with Conversational Language Understanding, you create an app; and then add entities, intents, and utterances to define the commands you want the app to execute.

1. In a new browser tab, open the Language Studio portal at <https://language.azure.com> and sign in using the Microsoft account associated with your Azure subscription.
2. If prompted to choose a Language resource, select the following settings:
 - **Azure directory:** The Azure directory containing your subscription.
 - **Azure subscription:** Your Azure subscription.
 - **Language resource:** The Language resource you created previously.

! **Tip** If you are *not* prompted to choose a language resource, it may be because you have multiple Language resources in your subscription; in which case:

- On the bar at the top of the page, click the **Settings** (⚙️) button.
- On the **Settings** page, view the **Resources** tab.
- Select your language resource, and click **Switch resource**.
- At the top of the page, click **Language Studio** to return to the Language Studio home page.

- At the top of the portal, in the **Create new** menu, select **Conversational language understanding**.
- In the **Create a project** dialog box, on the **Enter basic information** page, enter the following details and click **Next**:

- **Name:** *Create a unique name*
- **Description:** Simple home automation
- **Utterances primary language:** English
- **Enable multiple languages in project:** *Do not select*

Create a project

● Enter basic information

○ Review and finish

Enter basic information

Enter the basic information for your conversation project such as name and description and select the language which your utterances are in.

Azure resource

[Change Azure resource](#)

Name *

Description

Simple home automation

Utterances primary language *

English (US) ▾

☐ Enable multiple languages in project?

ⓘ

Next

Cancel

! **Tip** Write down your *project name*, you will use it later.

- On the *Review and finish* page, click **Create**.

Create intents, utterances, and entities

An *intent* is an action you want to perform - for example, you might want to switch on a light, or turn off a fan. In this case, you'll define two intents: one to switch on a device, and another to switch off a device. For each intent, you'll specify sample *utterances* that indicate the kind of language used to indicate the intent.

- In the **Schema definition** pane, ensure that **Intents** is selected Then click **Add**, and add an intent with the name **switch_on** (in lower-case) and click **Add intent**.

Language Studio

Projects

homeAutomation

Schema definition

Data labeling

Training jobs

Model performance

Schema definition

Add intents and entities to your schema.
user's intent.

Intents

Entities

+ Add

Delete

Intents

None

Add an intent

An intent represents a task or action that the user wants to perform. It is a purpose or goal expressed in a user's utterance.

For example, these utterances:

- "Book me a flight to Rio next week"
- "Fly me to Rio on the 24th"

would fall under "BookFlight" intent.

Name an intent

Intent name *

switch_on

Add intent

Cancel

"[] [] [] [] []"

+

→

✓

2. Select the **switch_on** intent. It will take you to the **Data labeling** page. In the **Intent** drop down, select **switch_on**. Next to the **switch_on** intent, type the utterance **turn the light on** and press **Enter** to submit this utterance to the list.

Language Studio

Projects

homeAutomation

Schema definition

Data labeling

Training jobs

Model performance

Deploying a model

Testing deployments

Project settings

Data labeling ✓ Saved

To add an utterance, first select an intent that your user intends to accomplish. Then provide an example of what they'd say in that situation and label the information to extract as entities. After labeling the utterances, you'll be ready to create a model with this data in [Training](#). For example, for an intent of Order Pizza, an utterance might be "I'd like to get a large pizza." The word "large" might then be labeled as a Pizza Size entity.

Training set

Testing set

Save changes

Upload utterance file

Move to

Edit

Delete

Filter

Filtered by: Intent: switch_on

We'll use these utterances to create your conversation model during training. A separate set of utterances can test the performance of your model. [Learn more about splitting data between training and testing sets.](#)

Intent

Utterance

switch_on

turn the light on

3. The language service needs at least five different utterance examples for each intent to sufficiently train the language model. Add five more utterance examples for the **switch_on** intent:

- switch on the fan
- put the fan on

- **put the light on**
- **switch on the light**
- **turn the fan on**

4. On the **Labeling entities for training** pane on the right-hand side of the screen, select **Labels**, then select **Add entity**. Type **device** (in lower-case), select **List** and select **Add entity**.

The screenshot shows the Azure Language Studio interface. On the left is a sidebar with navigation options: Language Studio, Projects, homeAutomation, Schema definition, Data labeling (selected), Training jobs, Model performance, Deploying a model, Testing deployments, and Project settings. The main area is titled 'Data labeling' and shows a 'Training set' tab. It includes instructions on how to add utterances and a list of utterances filtered by the intent 'switch_on'. On the right, the 'Labeling entities for training' pane is visible, showing the 'Labels' tab with an 'Add entity' button highlighted by an orange box. Below the button, it says 'You haven't created any entities.'

Add an entity

Name an entity

Entity name *

device

An entity can be extracted by different methods. They can be learned through context, matched from a list, or detected by a prebuilt. Every entity in your project is composed by one or more of these methods that are defined as your entity's components.

Learned **List** Prebuilt

The list component represents a fixed, closed set of related words along with their synonyms. The component performs an exact text match against the list of values you provide.

Book 2 business tickets to **Arrabury Airport**
airport

Add entity

Cancel

5. In the **turn the fan on** utterance, highlight the word "fan". Then in the list that appears, in the *Search for an entity* box select **device**.

Filtered by: Intent: **switch_on** ×

We'll use these utterances to create your conversation model during training. A separate set of utterances can test the performance of your model. [Learn more about splitting data between training and testing sets.](#)

○ Intent ↑

Utterance

switch_on

Write your example utterance and press enter

0/500

○ switch_on

turn the fan on

Search for an entity

device

View in tagging pane

○ switch_on

switch on

○ switch_on

switch on the light

6. Do the same for all the utterances. Label the rest of the *fan* or *light* utterances with the **device** entity. When you're finished, verify that you have the following utterances and make sure to select **Save changes**:

intent	utterance	entity
switch_on	Put on the fan	Device - <i>select fan</i>
switch_on	Put on the light	Device - <i>select light</i>
switch_on	Switch on the light	Device - <i>select light</i>
switch_on	Turn the fan on	Device - <i>select fan</i>
switch_on	Switch on the fan	Device - <i>select fan</i>
switch_on	Turn the light on	Device - <i>select light</i>

Language Studio

Projects

homeAutomation

Schema definition

Data labeling

Training jobs

Model performance

Data labeling

You have unsaved changes.

To add an utterance, first select an intent that your user intends to accomplish. Then provide an example of what they'd say in that situation and label the information to extract as entities. After labeling the utterances, you'll be ready to create a model with this data in [Training](#).

For example, for an intent of Order Pizza, an utterance might be "I'd like to get a large pizza." The word "large" might then be labeled as a Pizza Size entity.

Training set

Testing set

Save changes

Upload utterance file

Move to

Filter

7. In the pane on the left, click **Schema definition** and verify that your **switch_on** intent is listed. Then click **Add** and add a new intent with the name **switch_off** (in lower-case).

Language Studio

Projects

Schema definition

Data labeling

Training jobs

Model performance

Deploying a model

Schema definition

Add intents and entities to your schema. Intents are and can be extracted to help fulfill the user’s intent.

Intents

Entities

+ Add

Delete

Intents ↑ ↓

None

switch_on

8. Click on the **switch_off** intent. It will take you to the **Data labeling** page. In the **Intent** drop down, select **switch_off**. Next to the **switch_off** intent, add the utterance *turn the light off*.
9. Add five more utterance examples to the **switch_off** intent.
 - *switch off the fan*
 - *put the fan off*
 - *put the light off*
 - *turn off the light*
 - *switch the fan off*

10. Label the words *light* or *fan* with the **device** entity. When you’re finished, verify that you have the following utterances and make sure to select **Save changes**:

intent	utterance	entity
switch_off	Put the fan off	Device - <i>select fan</i>
switch_off	Put the light off	Device - <i>select light</i>
switch_off	Turn off the light	Device - <i>select light</i>
switch_off	Switch the fan off	Device - <i>select fan</i>
switch_off	Switch off the fan	Device - <i>select fan</i>
switch_off	Turn the light off	Device - <i>select light</i>

Train the model

Now you’re ready to use the intents and entities you have defined to train the conversational language model for your app.

1. On the left hand side of Language Studio, select **Training jobs**, then select **Start a training job**. Use the following settings:
 - **Train a new model:** *Selected and choose a model name*
 - **Training mode:** Standard training (free)
 - **Data Splitting:** *select Automatically split the testing set from the training data, keep default percentages*
 - Click **Train** at the bottom of the page.
2. Wait for training to complete.

Deploy and test the model

To use your trained model in a client application, you must deploy it as an endpoint to which the client applications can send new utterances; from which intents and entities will be predicted.

1. On the left-hand side of Language Studio, click **Deploying a model**.
2. Select your model name and click **Add deployment**. Use these settings:
 - **Create or select an existing deployment name:** *Select create a new deployment name. Add a unique name.*
 - **Assign trained model to your deployment name:** *Select the name of the trained model.*
 - Click **Deploy**

! **Tip** Write down your *deployment name*, you will use it later.

3. When the model is deployed, click **Testing deployments** on the left-hand side of the page, and then select your deployed model under **Deployment name**.
4. Enter the following text, and then select **Run the test**:

switch the light on

The screenshot shows the Language Studio interface. On the left, a sidebar lists various tools: Language Studio, Projects, homeAutomation, Schema definition, Data labeling, Training jobs, Model performance, Deploying a model, Testing deployments (selected), and Project settings. The main area is titled 'Testing deployments' and contains a 'Run the test' button. Below this is a 'Deployment name' dropdown menu. A text input field is labeled 'Enter your own text, or upload a text document' and contains the text 'switch the light on'. Below the input field, there are two tabs: 'Result' (selected) and 'JSON'. The 'Result' tab displays the predicted intent and entity. The intent is 'switch_on' with a confidence of 93.18%. The entity is 'light' with a confidence of 100.00%.

Intent
Top intent switch_on Confidence: 93.18%


Entities
device light Confidence: 100.00%

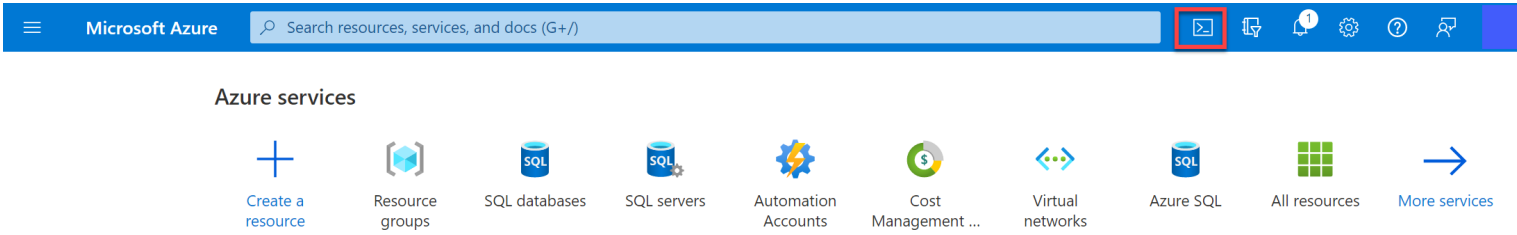
Review the result that is returned, noting that it includes the predicted intent (which should be **switch_on**) and the predicted entity (**device**) with confidence scores that indicates the probability the model calculated for the predicted intent and entity. The JSON tab shows the comparative confidence for each potential intent (the one with the highest confidence score is the predicted intent)

5. Clear the text box and test the model with the following utterances under *Enter your own text, or upload a text document*:
 - *turn off the fan*
 - *put the light on*
 - *put the fan off*

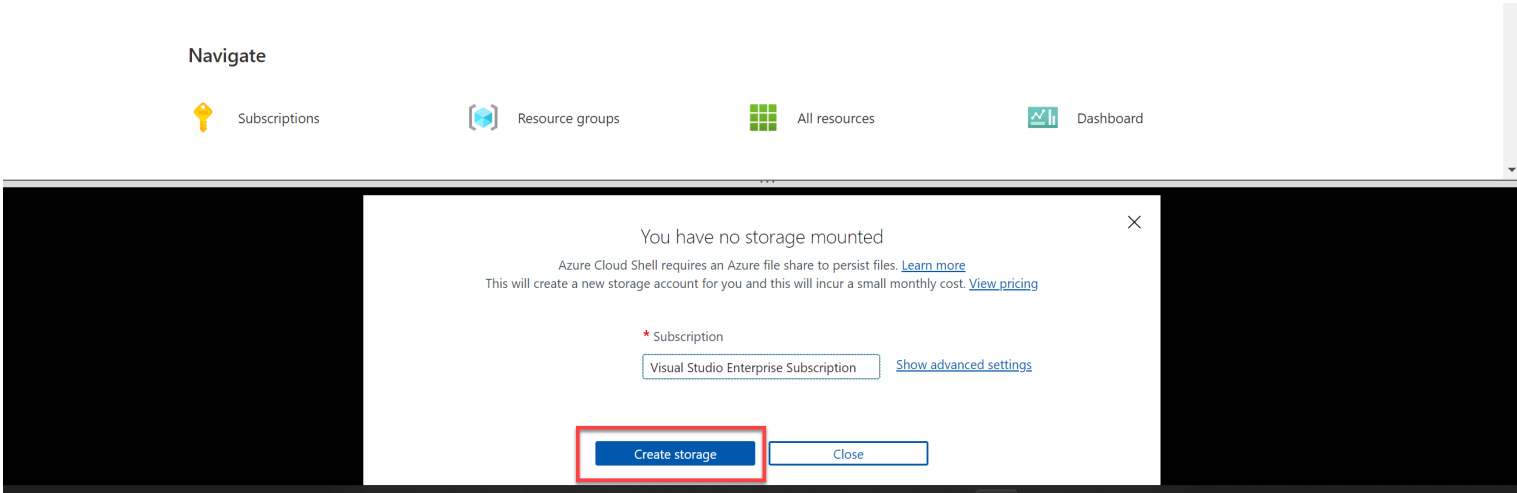
Run Cloud Shell

Now let's try out your deployed model. To do so, we'll use a command-line application that runs in the Cloud Shell on Azure.

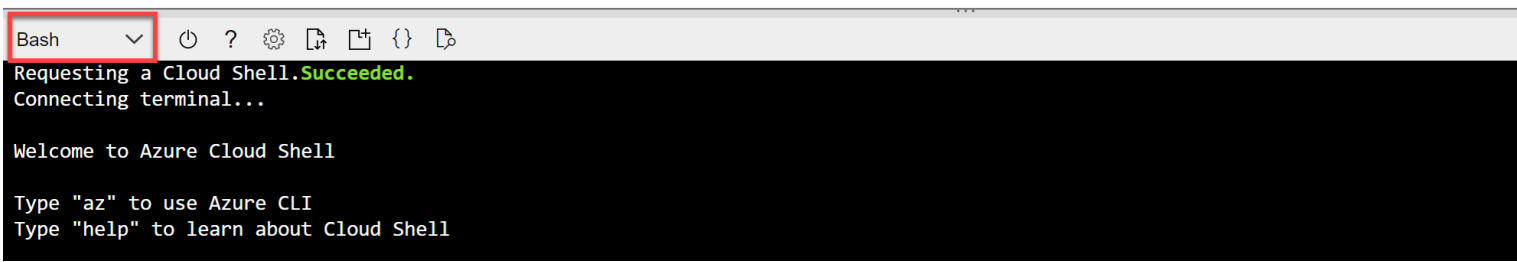
1. Leaving the browser tab with Language Studio open, switch back to browser tab containing the Azure portal.
2. In the Azure portal, select the  (Cloud Shell) button at the top of the page to the right of the search box. Clicking the button opens a Cloud Shell pane at the bottom of the portal.



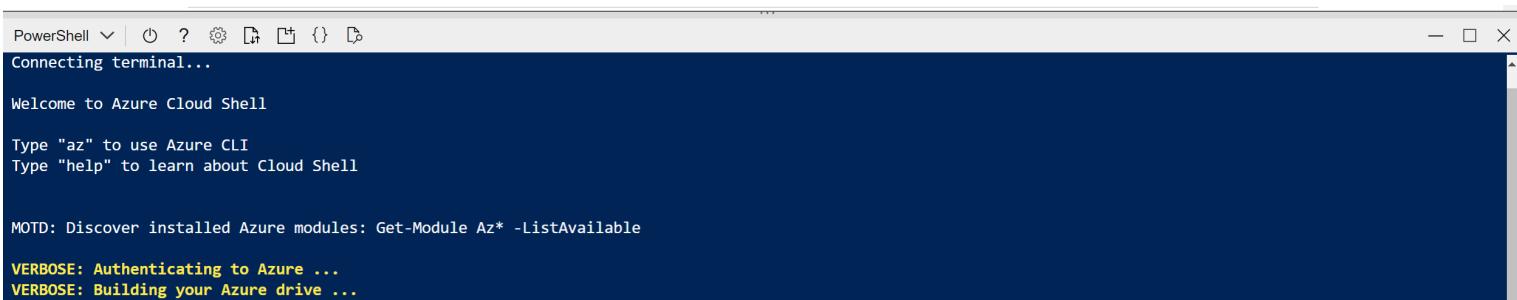
3. The first time you open the Cloud Shell, you may be prompted to choose the type of shell you want to use (Bash or PowerShell). Select **PowerShell**. If you do not see this option, skip the step.
4. If you are prompted to create storage for your Cloud Shell, ensure your subscription is specified and select **Create storage**. Then wait a minute or so for the storage to be created.



5. Make sure the type of shell indicated on the top left of the Cloud Shell pane is switched to *PowerShell*. If it is *Bash*, switch to *PowerShell* by using the drop-down menu.



6. Wait for PowerShell to start. You should see the following screen in the Azure portal:



Configure and run a client application

Now let's open and edit a pre-written script, which will run the client application.

1. In the command shell, enter the following command to download the sample application and save it to a folder called ai-900.

Code Copy

```
git clone https://github.com/MicrosoftLearning/AI-900-AIFundamentals ai-900
```


Note If you already used this command in another lab to clone the *ai-900* repository, you can skip this step.

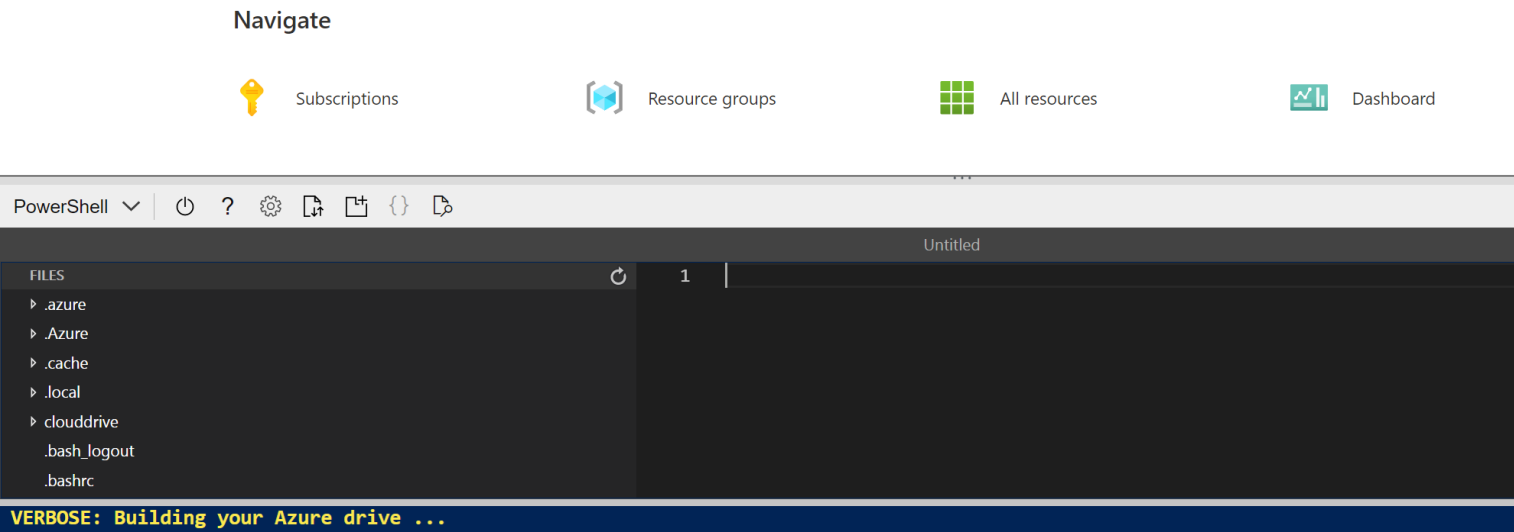
2. The files are downloaded to a folder named **ai-900**. Now we want to see all of the files in this folder and work with them. Type the following commands into the shell:

Code

Copy

```
cd ai-900
code .
```

Notice how the script opens up an editor like the one in the image below:



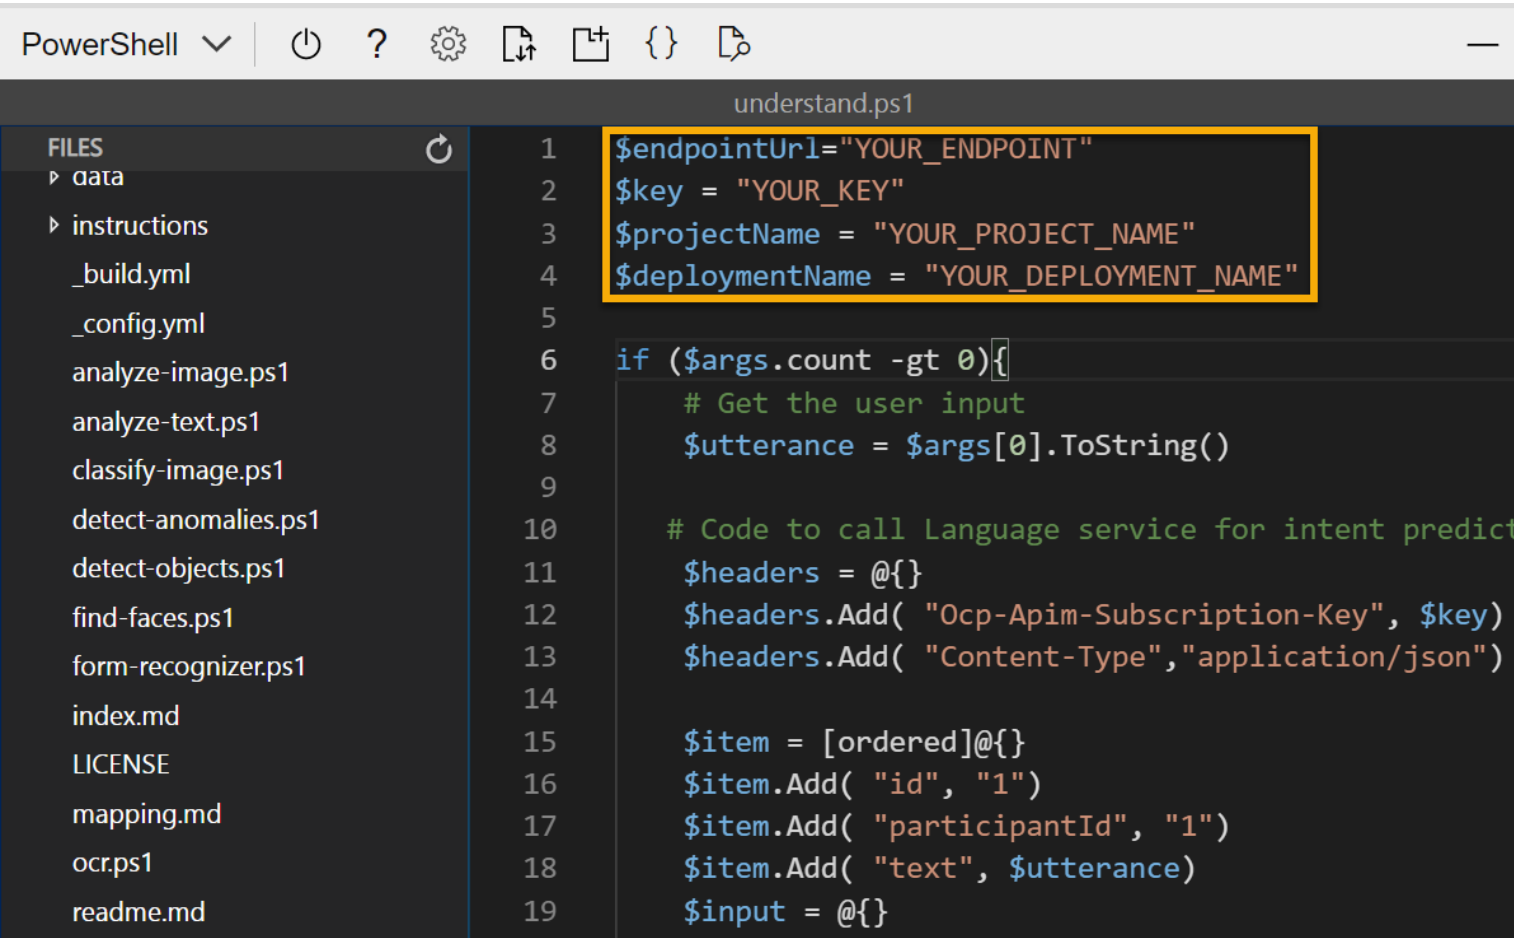
[Create a Language service resource](#)

[Run Cloud Shell](#)

[Configure and run a client application](#)

[Learn more](#)

3. In the **Files** pane on the left, select the **understand.ps1** file in the **ai-900** folder. This file contains some code that uses your Conversational Language Understanding model.




Don't worry too much about the details of the code. The important thing is that you'll use the instructions below to modify the file to specify the language model you trained.

4. Switch back to the browser tab containing **Language Studio**. Then in Language Studio, open the **Deploying a model** page and select your model. Then click the **Get prediction URL** button. The two pieces of information you need are in this dialog box:
- The endpoint for your model - you can copy the endpoint from the **Prediction URL** box.
 - The key for your model - the key is in the **Sample request** as the value for the **Ocp-Apim-Subscription-Key** parameter, and looks similar to **0ab1c23de4f56gh7i8901234jkl567m8**.
5. Copy the endpoint value, then switch back to the browser tab containing the Cloud Shell and paste it into the code editor, replacing **YOUR_ENDPOINT** (within the quotation marks). The repeat that process for the key, replacing **YOUR_KEY**.

6. Next, replace **YOUR_PROJECT_NAME** with the name of your project, and replace **YOUR_DEPLOYMENT_NAME** with the name of your deployed model. The first lines of code should look similar to what you see below:

Code	 Copy
<pre>\$endpointUrl="https://some-name.cognitiveservices.azure.com/language/..." \$key = "0ab1c23de4f56gh7i8901234jkl567m8" \$projectName = "name" \$deploymentName = "name"</pre>	


7. At the top right of the editor pane, use the ... button to open the menu and select **Save** to save your changes. Then open the menu again and select **Close Editor**.
8. In the PowerShell pane, enter the following command to run the code:

Code	 Copy
<pre>./understand.ps1 "Turn on the light"</pre>	

9. Review the results. The app should have predicted that the intended action is to switch on the light.
10. Now try another command:

Code	 Copy
<pre>./understand.ps1 "Switch the fan off"</pre>	

11. Review the results from this command. The app should have predicted that the intended action is to switch off the fan.
12. Experiment with a few more commands; including commands that the model was not trained to support, such as “Hello” or “switch on the oven”. The app should generally understand commands for which its language model is defined, and fail gracefully for other input.

 **Note** Each time you will need to start with **./understand.ps1** followed by the phrase. Include quotation marks around your phrase.

Learn more

This app shows only some of the capabilities of the Conversational Language Understanding feature of the Language service. To learn more about what you can do with this service, see the [Conversational Language Understanding page](#).